

Applied survival analysis

Constantin T Yiannoutsos, Ph.D.

January 13, 2016

1 Estimating the survival function

- The Kaplan Meier estimator
 - Review
 - Calculating the Kaplan-Meier estimator
 - Variance of the Kaplan-Meier estimator
- The Life Table estimator of survival
 - Notation
 - Calculating the life table estimator
 - Computer implementation

2 Estimates of the cumulative hazard

- The Nelson-Aalen estimator
- The Fleming-Harrington estimator

Estimating the Survival Function

We will consider three methods for estimating a survivorship function

$$S(t) = Pr(T \geq t)$$

without resorting to parametric methods:

- (1) **Kaplan-Meier**
- (2) **Life-table** (Actuarial Estimator)
- (3) **Cumulative hazard estimator**

The Kaplan-Meier estimator

The Kaplan-Meier (or KM) estimator is probably the most popular approach. It can be justified from several perspectives:

- product limit estimator
- likelihood justification
- redistribute to the right estimator

We will start with an intuitive motivation based on conditional probabilities, then review some of the other justifications.

Motivation

First, consider an example where there is no censoring.

The following are times of remission (weeks) for 21 leukemia patients receiving control treatment (Table 1.1 of Cox & Oakes):

1, 1, 2, 2, 3, 4, 4, 5, 5, 8, 8, 8, 8, 11, 11, 12, 12, 15, 17, 22, 23

How would we estimate $S(10)$, the probability that an individual survives to time 10 or later?

What about $\tilde{S}(8)$? Is it $\frac{12}{21}$ or $\frac{8}{21}$?

Let's construct a table of $\hat{S}(t)$:

Values of t	$\hat{S}(t)$
$t \leq 1$	$21/21=1.000$
$1 < t \leq 2$	$19/21=0.905$
$2 < t \leq 3$	$17/21=0.809$
$3 < t \leq 4$	
$4 < t \leq 5$	
$5 < t \leq 8$	
$8 < t \leq 11$	
$11 < t \leq 12$	
$12 < t \leq 15$	
$15 < t \leq 17$	
$17 < t \leq 22$	
$22 < t \leq 23$	


The Empirical Survival Function

When there is no censoring, the general formula is:

$$\tilde{S}(t) = \frac{\# \text{ individuals with } T \geq t}{\text{total sample size}}$$

Example for the leukemia data (control arm)

Bone Marrow Transplant for Non-Hodgkin's Lymphoma



When there is censoring

Consider the treated group from Table 1.1 of Cox and Oakes:

$6^+, 6, 6, 6, 7, 9^+, 10^+, 10, 11^+, 13, 16, 17^+$
 $19^+, 20^+, 22, 23, 25^+, 32^+, 32^+, 34^+, 35^+$

[Note: times with $^+$ are right censored]

We know $S(6) = 21/21$, because everyone survived at least until time 6 or greater. But, we can't say $S(7) = 17/21$, because we don't know the status of the person who was censored at time 6.

In a 1958 paper in the *Journal of the American Statistical Association*, Kaplan and Meier proposed a way to nonparametrically estimate $S(t)$, even in the presence of censoring. The method is based on the ideas of **conditional probability**.

A quick review of conditional probability

Suppose A and B are two events.

Then,

$$P(A|B) = \frac{P(A \cap B)}{P(B)}$$

Multiplication law of probability: can be obtained from the above relationship, by multiplying both sides by $P(B)$:

$$P(A \cap B) = P(A|B) P(B)$$

Extension to more than 2 events

Suppose $A_1, A_2 \dots A_k$ are k different events. Then, the probability of all k events happening together can be written as a product of conditional probabilities:

$$\begin{aligned} P(A_1 \cap A_2 \dots \cap A_k) &= P(A_k | A_{k-1} \cap \dots \cap A_1) \times \\ &\quad \times P(A_{k-1} | A_{k-2} \cap \dots \cap A_1) \\ &\quad \dots \\ &\quad \times P(A_2 | A_1) \\ &\quad \times P(A_1) \end{aligned}$$

Bringing it all together

Now, let's apply these ideas to estimate $S(t)$:

Suppose $a_k < t \leq a_{k+1}$. Then

$$\begin{aligned} S(t) &= P(T \geq a_{k+1}) \\ &= P(T \geq a_1, T \geq a_2, \dots, T \geq a_{k+1}) \\ &= P(T \geq a_1) \times \prod_{j=1}^k P(T \geq a_{j+1} | T \geq a_j) \\ &= \prod_{j=1}^k [1 - P(T = a_j | T \geq a_j)] \\ &= \prod_{j=1}^k [1 - \lambda_j] \end{aligned}$$

So,

$$\begin{aligned}\hat{S}(t) &\cong \prod_{j=1}^k \left(1 - \frac{d_j}{r_j}\right) \\ &= \prod_{j:a_j < t} \left(1 - \frac{d_j}{r_j}\right)\end{aligned}$$

d_j is the number of deaths at a_j

r_j is the number at risk at a_j

Intuition behind the Kaplan-Meier Estimator

Think of dividing the observed timespan of the study into a series of fine intervals so that there is a separate interval for each time of death or censoring:



Using the law of conditional probability,

$$Pr(T \geq t) = \prod_j Pr(\text{survive } j\text{-th interval } I_j \mid \text{survived to start of } I_j)$$

where the product is taken over all the intervals including or preceding time t .

What can happen

These are possibilities for each interval:

- (1) **No events (death or censoring)** - conditional probability of surviving the interval is 1
- (2) **Censoring** - assume they survive to the end of the interval, so that the conditional probability of surviving the interval is 1
- (3) **Death, but no censoring** - conditional probability of *not* surviving the interval is $\# \text{ deaths } (d) \text{ divided by } \# \text{ 'at risk' } (r) \text{ at the beginning of the interval}$. So the conditional probability of surviving the interval is $1 - (d/r)$.
- (4) **Tied deaths and censoring** - assume censorings last to the end of the interval, so that conditional probability of surviving the interval is still $1 - (d/r)$

General Formula for j th interval

It turns out we can write a general formula for the conditional probability of surviving the j -th interval that holds for all 4 cases:

$$1 - \frac{d_j}{r_j}$$

We could use the same approach by grouping the event times into intervals (say, one interval for each month), and then counting up the number of deaths (events) in each to estimate the probability of surviving the interval (this is called the *lifetable estimate*).

However, the assumption that those censored last until the end of the interval wouldn't be quite accurate, so we would end up with a cruder approximation.

The Kaplan-Meier - product-limit - estimator

As the intervals get finer and finer, the approximations made in estimating the probabilities of getting through each interval become smaller and smaller, so that the estimator converges to the true $S(t)$.

This intuition clarifies why an alternative name for the KM is the product limit estimator.

The Kaplan-Meier estimator of the survivorship function (or survival probability) $S(t) = Pr(T \geq t)$ is:

$$\hat{S}(t) = \prod_{j: \tau_j < t} \frac{r_j - d_j}{r_j} = \prod_{j: \tau_j < t} \left(1 - \frac{d_j}{r_j}\right)$$

where,

- τ_1, \dots, τ_K are the K distinct death times observed in the sample
- d_j is the number of deaths at τ_j
- r_j is the number of individuals “at risk” right before the j -th death time (everyone dead or censored at or after that time).
- c_j is the number of censored observations between the j -th and $(j+1)$ -st death times. Censorings tied at τ_j are included in c_j

Note: two useful formulas are:

$$(1) \quad r_j = r_{j-1} - d_{j-1} - c_{j-1}$$

$$(2) \quad r_j = \sum_{l \geq j} (c_l + d_l)$$

The Cox and Oakes example

Make a table with a row for every death or censoring time:


τ_j	d_j	c_j	r_j	$1 - (d_j/r_j)$	$\hat{S}(\tau_j^+)$
6	3	1	21	$\frac{18}{21} = 0.857$	
7	1	0	17		
9	0	1	16		
10					
11					
13					
16					
17					
19					
20					
22					
23					

Note that:

- $\hat{S}(t^+)$ only changes at death (failure) times
- $\hat{S}(t^+)$ is 1 up to the first death time
- $\hat{S}(t^+)$ only goes to 0 if the last event is a death

KM plot for treated leukemia patients

Bone Marrow Transplant for Non-Hodgkin's Lymphoma



Statistical software

Note: most statistical software packages summarize the KM survival function at τ_j^+ , i.e., *just after* the time of the j -th failure.

In other words, they provide $\hat{S}(\tau_j^+)$.

When there is no censoring, the empirical survival estimate would then be:

$$\tilde{S}(t^+) = \frac{\# \text{ individuals with } T > t}{\text{total sample size}}$$

Computer implementation

The output from R is as follows:

time	n.risk	n.event	survival	std.err	lower 95% CI	upper 95% CI
1	12	1	0.9167	0.0798	0.53898	0.988
2	11	3	0.6667	0.1361	0.33702	0.860
3	8	1	0.5833	0.1423	0.27014	0.801
5	7	1	0.5000	0.1443	0.20848	0.736
6	6	1	0.4167	0.1423	0.15247	0.665
7	5	1	0.3333	0.1361	0.10270	0.588
8	4	1	0.2500	0.1250	0.06014	0.505
16	3	1	0.1667	0.1076	0.02651	0.413
17	2	1	0.0833	0.0798	0.00505	0.311
34	1	1	0.0000	NaN	NA	NA

Two other justifications for the KM estimator

I. Likelihood-based derivation (Cox and Oakes)

For a discrete failure time variable, define:

- d_j number of failures at a_j
- r_j number of individuals at risk at a_j
(including those censored at a_j).
- λ_j Pr(death) in j -th interval
(conditional on survival to start of interval)

The likelihood is that of g independent binomials:

$$L(\lambda) = \prod_{j=1}^g \lambda_j^{d_j} (1 - \lambda_j)^{r_j - d_j}$$

Likelihood justification (cont'd)

Therefore, the **maximum likelihood estimator** of λ_j is:

$$\hat{\lambda}_j = d_j / r_j$$

Now we plug in the MLE's of λ to estimate $S(t)$:

$$\begin{aligned}\hat{S}(t) &= \prod_{j:a_j < t} (1 - \hat{\lambda}_j) \\ &= \prod_{j:a_j < t} \left(1 - \frac{d_j}{r_j}\right)\end{aligned}$$

The “Redistribute to the right” justification

II. Redistribute to the right justification (Efron, 1967)

In the absence of censoring, $\hat{S}(t)$ is just the proportion of individuals with $T \geq t$. The idea behind Efron's approach is to spread the contributions of censored observations out over all the possible times to their right.

Algorithm

The redistribute to the right algorithm is as follows:

- Step (1): arrange the n observed times (deaths or censorings) in increasing order. If there are ties, put censored after deaths.
- Step (2): Assign weight $(1/n)$ to each time.
- Step (3): Moving from left to right, each time you encounter a censored observation, distribute its mass to all times to its right.
- Step (4): Calculate \hat{S}_j by subtracting the final weight for time j from \hat{S}_{j-1}

Example of “redistribute to the right” algorithm

Consider the following event times:

2, 2.5+, 3, 3, 4, 4.5+, 5, 6, 7

The algorithm goes as follows:

(Step 1) Times	Step 2	Step 3a	Step 3b	(Step 4) $\hat{S}(\tau_j)$
2	1/9=0.11			0.889
2.5+	1/9=0.11	0		0.889
3	2/9=0.22	0.25		0.635
4	1/9=0.11	0.13		0.508
4.5+	1/9=0.11	0.13	0	0.508
5	1/9=0.11	0.13	0.17	0.339
6	1/9=0.11	0.13	0.17	0.169
7	1/9=0.11	0.13	0.17	0.000

This comes out the same as the product limit approach.

Properties of the KM estimator

In the case of no censoring:

$$\hat{S}(t) = \tilde{S}(t) = \frac{\# \text{ deaths at } t \text{ or greater}}{n}$$

where n is the number of individuals in the study.

This is just like an estimated probability from a binomial distribution, so we have:

$$\hat{S}(t) \simeq \mathcal{N}(S(t), S(t)[1 - S(t)]/n)$$

How does censoring affect this?

- $\hat{S}(t)$ is still approximately normal
- The mean of $\hat{S}(t)$ converges to the true $S(t)$
- The variance is a bit more complicated (since the denominator n includes some censored observations).

Once we get the variance, then we can construct (pointwise) $(1 - \alpha)\%$ confidence bands about $\hat{S}(t)$:

$$\hat{S}(t) \pm z_{1-\alpha/2} \text{se}[\hat{S}(t)]$$

Greenwood's formula (Collett 2.1.3)

We can think of the KM estimator as

$$\hat{S}(t) = \prod_{j: \tau_j < t} (1 - \hat{\lambda}_j)$$

where $\hat{\lambda}_j = d_j/r_j$. Since the $\hat{\lambda}_j$'s are just binomial proportions, we can apply standard likelihood theory to show that each $\hat{\lambda}_j$ is approximately normal, with mean the true λ_j , and

$$\text{var}(\hat{\lambda}_j) \approx \frac{\hat{\lambda}_j(1 - \hat{\lambda}_j)}{r_j}$$

The $\hat{\lambda}_j$'s are independent in large samples. Since $\hat{S}(t)$ is a function of the λ_j 's, we can estimate its variance using the **delta method**:

If Y is normal with mean μ and variance σ^2 , then $g(Y)$ is approximately normally distributed with mean $g(\mu)$ and variance $[g'(\mu)]^2 \sigma^2$.

Two specific examples of the delta method

(A) $Z = \log(Y)$

$$\text{then } Z \sim N \left[\log(\mu), \left(\frac{1}{\mu} \right)^2 \sigma^2 \right]$$

(B) $Z = \exp(Y)$

$$\text{then } Z \sim N [e^\mu, [e^\mu]^2 \sigma^2]$$

The examples above use the following results from calculus:

$$\frac{d}{dx} \log u = \frac{1}{u} \left(\frac{du}{dx} \right)$$

$$\frac{d}{dx} e^u = e^u \left(\frac{du}{dx} \right)$$

Greenwood's formula (continued)

Instead of dealing with $\hat{S}(t)$ directly, we will look at its log:

$$\log[\hat{S}(t)] = \sum_{j:\tau_j < t} \log(1 - \hat{\lambda}_j)$$

Thus, by approximate independence of the $\hat{\lambda}_j$'s,

$$\text{var}(\log[\hat{S}(t)]) = \sum_{j:\tau_j < t} \text{var}[\log(1 - \hat{\lambda}_j)]$$

By (A)

$$\begin{aligned} \text{var}(\log[\hat{S}(t)]) &= \sum_{j:\tau_j < t} \left(\frac{1}{1 - \hat{\lambda}_j} \right)^2 \text{var}(\hat{\lambda}_j) \\ &= \sum_{j:\tau_j < t} \left(\frac{1}{1 - \hat{\lambda}_j} \right)^2 \hat{\lambda}_j(1 - \hat{\lambda}_j)/r_j \\ &= \sum_{j:\tau_j < t} \frac{\hat{\lambda}_j}{(1 - \hat{\lambda}_j)r_j} = \sum_{j:\tau_j < t} \frac{d_j}{(r_j - d_j)r_j} \end{aligned}$$

Back to confidence intervals

For a 95% confidence interval, we could use

$$\hat{S}(t) \pm z_{1-\alpha/2} \text{se}[\hat{S}(t)]$$

where $\text{se}[\hat{S}(t)]$ is calculated using Greenwood's formula.

Problem: This approach can yield values > 1 or < 0 .

Better approach: Get a 95% confidence interval for

$$L(t) = \log(-\log(S(t)))$$

Since this quantity is unrestricted, the confidence interval will be in the right range when we transform back.

Why does this work?

To see why this works, note the following:

- Since $\hat{S}(t)$ is an estimated probability

$$0 \leq \hat{S}(t) \leq 1$$

- Taking the log of $\hat{S}(t)$ and bounds:

$$-\infty \leq \log[\hat{S}(t)] \leq 0$$

- Taking the opposite:

$$0 \leq -\log[\hat{S}(t)] \leq \infty$$

- Taking the log again:

$$-\infty \leq \log \left[-\log[\hat{S}(t)] \right] \leq \infty$$

To transform back, reverse steps with $S(t) = \exp(-\exp(L(t)))$

Log-log Approach for Confidence Intervals

To calculate the log-log transformed confidence interval we proceed as follows:

- (1) Define $L(t) = \log(-\log(S(t)))$
- (2) Form a 95% confidence interval for $L(t)$ based on $\hat{L}(t)$, yielding $[\hat{L}(t) - A, \hat{L}(t) + A]$
- (3) Since $S(t) = \exp(-\exp(L(t)))$, the confidence bounds for the 95% CI on $S(t)$ are:

$$[\exp(-e^{(\hat{L}(t)+A)}), \exp(-e^{(\hat{L}(t)-A)})]$$

(note that the upper and lower bounds switch)

- (4) Substituting $\hat{L}(t) = \log(-\log(\hat{S}(t)))$ back into the above bounds, we get confidence bounds of

$$([\hat{S}(t)]^{e^A}, [\hat{S}(t)]^{e^{-A}})$$

What is A?

- A is $1.96 \text{ se}(\hat{L}(t))$
- To calculate this, we need to calculate

$$\text{var}(\hat{L}(t)) = \text{var} \left[\log(-\log(\hat{S}(t))) \right]$$

- From our previous calculations, we know

$$\text{var}(\log[\hat{S}(t)]) = \sum_{j:\tau_j < t} \frac{d_j}{(r_j - d_j)r_j}$$

Applying the delta method

Applying the Delta method as in example (A), we get:

$$\begin{aligned}\text{var}(\hat{L}(t)) &= \text{var}(\log(-\log[\hat{S}(t)])) \\ &= \frac{1}{[\log \hat{S}(t)]^2} \sum_{j:\tau_j < t} \frac{d_j}{(r_j - d_j)r_j}\end{aligned}$$

- We take the square root of the above to get $se(\hat{L}(t))$, and then form the confidence intervals as:

$$\hat{S}(t)e^{\pm 1.96 \, se(\hat{L}(t))}$$

- This is the approach that Stata uses. Splus also gives an option to calculate these bounds.

Summary of confidence intervals on $S(t)$

To obtain confidence intervals for $S(t)$ we proceed as follows: We calculate $\hat{S}(t) \pm 1.96 \text{se}[\hat{S}(t)]$ where $\text{se}[\hat{S}(t)]$ is calculated using Greenwood's formula, and replace negative lower bounds by 0 and upper bounds greater than 1 by 1 (not very satisfactory).

- Recommended by Collett
- This is the default using SAS

Confidence intervals for $S(t)$ (con'd)

We use a log transformation to stabilize the variance and allow for non-symmetric confidence intervals. This is what is normally done for the confidence interval of an estimated odds ratio.

- Use $\text{var}[\log(\hat{S}(t))] = \sum_{j:\tau_j < t} \frac{d_j}{(r_j - d_j)r_j}$ already calculated as part of Greenwood's formula
- This is the default in R for example

Confidence intervals for $S(t)$ (con'd)

We now use the log-log transformation just described

- Somewhat complicated, but always yields proper bounds
- This is the default in Stata!

Software for Kaplan-Meier curves

- R - `surv.fit(time,censor)`
- Stata - `stset` and `sts` commands
- SAS - `PROC LIFETEST`

Defaults for confidence interval calculations

- Splus - “log” $\Rightarrow \log S(t) \pm 1.96 \text{ se}[\log(\hat{S}(t))]$
but Splus will also give either of the other two options if you request them.
- Stata - “log-log” $\Rightarrow \hat{L}(t) \pm 1.96 \text{ se}[\hat{L}(t)]$
where $L(t) = \log[-\log(S(t))]$
- SAS - “plain” $\Rightarrow \hat{S}(t) \pm 1.96 \text{ se}[\hat{S}(t)]$

R output for treated leukemia patients

time	n.risk	n.event	survival	std.err	lower 95% CI	upper 95% CI
1	12	1	0.917	0.0798	0.5390	0.988
2	11	2	0.750	0.1250	0.4084	0.912
3	8	1	0.656	0.1402	0.3204	0.856
5	7	1	0.562	0.1482	0.2437	0.791
6	6	1	0.469	0.1503	0.1762	0.718
7	5	0	0.469	0.1503	0.1762	0.718
8	4	1	0.352	0.1517	0.0956	0.628
16	3	0	0.352	0.1517	0.0956	0.628
17	2	1	0.176	0.1456	0.0120	0.505
34	1	0	0.176	0.1456	0.0120	0.505

Summary statistics

n	events	median	0.95LCL	0.95UCL
12	8	6	2	NA

The Life Table estimator of survival

We said that we would consider the following three methods for estimating a survivorship function

$$S(t) = Pr(T \geq t)$$

without resorting to parametric methods:

- (1) ✓ **Kaplan-Meier**
- (2) \implies **Life-table** (Actuarial Estimator)
- (3) \implies **Cumulative hazard estimator**

The Life Table or Actuarial Estimator

The Life Table or Actuarial Estimator of survival:

- one of the oldest techniques around
- used by actuaries, demographers, etc.
- **applies when the data are grouped**

Our goal is still to estimate the survival function, hazard, and density function, but this is complicated by the fact that we don't know exactly when during each time interval an event occurs. Lee (section 4.2) provides a good description of lifetable methods, and distinguishes several types according to the data sources:

Population Life Tables

Here are some kinds of Life Tables:

- **cohort life table** - describes the mortality experience from birth to death for a particular cohort of people born at about the same time. People at risk at the start of the interval are those who survived the previous interval.
- **current life table** - constructed from (1) census information on the number of individuals alive at each age, for a given year and (2) vital statistics on the number of deaths or failures in a given year, by age. This type of lifetable is often reported in terms of a hypothetical cohort of 100,000 people.

Generally, censoring is not an issue for Population Life Tables.

Clinical Life tables

Clinical Life Tables apply to grouped survival data from studies in patients with specific diseases. Because patients can enter the study at different times, or be lost to follow-up, censoring must be allowed.

Notation

- the j -th time interval is $[t_{j-1}, t_j)$
- c_j - the number of censorings in the j -th interval
- d_j - the number of failures in the j -th interval
- r_j is the number entering the interval

Example: Angina pectoris

Example: 2418 Males with Angina Pectoris (Lee, p.91):

Year after Diagnosis	j	d_j	c_j	r_j	$r'_j = r_j - c_j/2$
[0, 1)	1	456	0	2418	2418.0
[1, 2)	2	226	39	1962	1942.5 (1962 - $\frac{39}{2}$)
[2, 3)	3	152	22	1697	1686.0
[3, 4)	4	171	23	1523	1511.5
[4, 5)	5	135	24	1329	1317.0
[5, 6)	6	125	107	1170	1116.5
[6, 7)	7	83	133	938	871.5
etc..					

Estimating the survivorship function

We could apply the K-M formula directly to the numbers in the table on the previous page, estimating $S(t)$ as

$$\hat{S}(t) = \prod_{j: \tau_j < t} \left(1 - \frac{d_j}{r_j}\right)$$

However, this approach is unsatisfactory for grouped data.... it treats the problem as though it were in discrete time, with events happening only at 1 yr, 2 yr, etc. In fact, what we are trying to calculate here is the conditional probability of dying within the interval, given survival to the beginning of it.

What should we do with the censored people?

We can assume that censorings occur:

- at the beginning of each interval: $r'_j = r_j - c_j$
- at the end of each interval: $r'_j = r_j$
- on average halfway through the interval:

$$r'_j = r_j - c_j/2$$

The last assumption yields the Actuarial Estimator. It is appropriate if censorings occur uniformly throughout the interval.

Constructing the Life Table

First, some additional notation for the j -th interval, $[t_{j-1}, t_j)$:

- **Midpoint** (t_{mj}) - useful for plotting the density and the hazard function
- **Width** ($b_j = t_j - t_{j-1}$) needed for calculating the hazard in the j -th interval

Quantities estimated:

- Conditional probability of dying is $\hat{q}_j = d_j / r'_j$
- Conditional probability of surviving is $\hat{p}_j = 1 - \hat{q}_j$
- Cumulative probability of surviving at t_j :

$$\hat{S}(t_j) = \prod_{\ell \leq j} \hat{p}_\ell = \prod_{\ell \leq j} \left(1 - \frac{d_\ell}{r'_\ell} \right)$$

Some important points to note

Here are some important points to note:

- Because the intervals are defined as $[t_{j-1}, t_j)$, the first interval typically starts with $t_0 = 0$.
- Stata estimates the survival function at the right-hand endpoint of each interval, i.e., $S(t_j)$
- However, SAS and R estimate the survival function at the left-hand endpoint, $S(t_{j-1})$.
- The implication in SAS and R is that $\hat{S}(t_0) = 1$ and $\hat{S}(t_1) = p_1$

Other quantities estimated at the midpoint of the j -th interval

Other quantities estimated are as follows:

- **Hazard** in the j -th interval:

$$\hat{\lambda}(t_{mj}) = \frac{d_j}{b_j(r'_j - d_j/2)} = \frac{\hat{q}_j}{b_j(1 - \hat{q}_j/2)}$$

the number of deaths in the interval divided by the average number of survivors at the midpoint

- **density** at the midpoint of the j -th interval:

$$\hat{f}(t_{mj}) = \frac{\hat{S}(t_{j-1}) - \hat{S}(t_j)}{b_j} = \frac{\hat{S}(t_{j-1}) \hat{q}_j}{b_j}$$

Note: Another way to get this is:

$$\hat{f}(t_{mj}) = \hat{\lambda}(t_{mj})\hat{S}(t_{mj}) = \hat{\lambda}(t_{mj})[\hat{S}(t_j) + \hat{S}(t_{j-1})]/2$$

Constructing the Life Table using R

Uses the `lifetab` command. The results are as follows:

	nsubs	nlost	nrisk	nevent	surv	pdf	hazard	se.surv	se.pdf	se.hazard
0-1	2418	0	2418.0	456	1.0000000	0.18858561	0.20821918	0.000000000	0.007955134	0.009697769
1-2	1962	39	1942.5	226	0.8114144	0.09440394	0.12353102	0.007955134	0.005975178	0.008201472
2-3	1697	22	1686.0	152	0.7170105	0.06464151	0.09440994	0.009179397	0.005069200	0.007649121
3-4	1523	23	1511.5	171	0.6523689	0.07380423	0.11991585	0.009734736	0.005428013	0.009153696
4-5	1329	24	1317.0	135	0.5785647	0.05930618	0.10804322	0.010138361	0.004945997	0.009285301
5-6	1170	107	1116.5	125	0.5192585	0.05813463	0.11859583	0.010304216	0.005033980	0.010588867
6-7	938	133	871.5	83	0.4611239	0.04391656	0.10000000	0.010379949	0.004690538	0.010962697
7-8	722	102	671.0	74	0.4172073	0.04601094	0.11671924	0.010450930	0.005175094	0.013545211
8-9	546	68	512.0	51	0.3711964	0.03697464	0.10483042	0.010578887	0.005024599	0.014659017
9-10	427	64	395.0	42	0.3342218	0.03553750	0.11229947	0.010717477	0.005307615	0.017300846
10-11	321	45	298.5	43	0.2986843	0.04302654	0.15523466	0.010890741	0.006269963	0.023601647
11-12	233	53	206.5	34	0.2556577	0.04209376	0.17941953	0.011124244	0.006847514	0.030646128
12-13	146	33	129.5	18	0.2135639	0.02968456	0.14937759	0.011396799	0.006682743	0.035110295
13-14	95	27	81.5	9	0.1838794	0.02030570	0.11688312	0.011765989	0.006514794	0.038894448
14-15	59	23	47.5	6	0.1635737	0.02066194	0.13483146	0.012259921	0.008035120	0.054919485
15-NA	30	30	15.0	0	0.1429117	NA	NA	0.013300258	NA	NA

The above also estimates the hazard function, $\hat{\lambda}_j$, and its standard error.

Calculating the Actuarial Life Table estimator when the data are not grouped


Suppose we wish to use the actuarial method, but the data do not come grouped.

Consider the treated nursing home patients, with length of stay (los) grouped into 100-day intervals:

	nsubs	nlost	nrisk	nevent	surv	pdf	hazard	se.s urv	se.pdf	se.hazard
0-1	710	0	710.0	328	1.0000000	-0.076995305	-0.10012210	0.00000000	-0.0031183857	0.005273042
1-2	382	0	382.0	86	0.5380282	0.121126761	0.25368732	0.01871031	0.0122448649	0.027134846
2-3	296	0	296.0	65	0.4169014	0.015258216	0.04111322	0.01850370	0.0018038390	0.005060530
3-4	231	0	231.0	38	0.3253521	-0.010704225	-0.03584906	0.01758273	-0.0016893472	0.005792082
4-5	193	1	192.5	32	0.2718310	0.007531248	0.03021719	0.01669692	0.0013007076	0.005319701
5-6	160	0	160.0	13	0.2266435	0.018414784	0.08469055	0.01571642	0.0050592654	0.023467865
6-7	147	0	147.0	13	0.2082287	0.018414784	0.09252669	0.01524675	0.0050592654	0.025634809
7-8	134	30	119.0	10	0.1898139	-0.002278679	-0.01253133	0.01472901	-0.0007119474	0.003958941
8-9	94	29	79.5	4	0.1738632	0.008747833	0.05161290	0.01432896	0.0043230020	0.025797857
9-10	61	30	46.0	4	0.1651153	0.014357856	0.09090909	0.01425996	0.0069708766	0.045407564
10-NA	27	27	13.5	0	0.1507575	NA	NA	0.01471649	NA	NA


Example: Duration of stay in nursing homes

Estimated Survival:



Example: Duration of stay in nursing homes

Estimated hazard:



Estimating the cumulative hazard: The Nelson-Aalen estimator

Suppose we want to estimate $\Lambda(t) = \int_0^t \lambda(u) du$, the cumulative hazard at time t .

Just as we did for the KM, think of dividing the observed timespan of the study into a series of fine intervals so that there is only one event per interval:



The cumulative hazard $\Lambda(t)$ can then be approximated by a sum:

$$\hat{\Lambda}(t) = \sum_j \lambda_j \Delta$$

where the sum is over intervals, λ_j is the value of the hazard in the j -th interval and Δ is the width of each interval. Since $\hat{\lambda}\Delta$ is approximately the probability of dying in the interval, we can further approximate by

$$\hat{\Lambda}(t) = \sum_j d_j / r_j$$

It follows that $\Lambda(t)$ will change only at death times, and hence we write the Nelson-Aalen estimator as:

$$\hat{\Lambda}_{NA}(t) = \sum_{j: \tau_j < t} d_j / r_j$$

The Fleming-Harrington (FH) estimator

			D		C		C	D	D	D
r_j	n	n	n	$n-1$	$n-1$	$n-2$	$n-2$	$n-3$	$n-4$	
d_j	0	0	1	0	0	0	0	1	1	
c_j	0	0	0	0	1	0	1	0	0	
$\hat{\lambda}(t_j)$	0	0	$1/n$	0	0	0	0	$\frac{1}{n-3}$	$\frac{1}{n-4}$	
$\hat{\Lambda}(t_j)$	0	0	$1/n$	$1/n$	$1/n$	$1/n$	$1/n$			

Once we have $\hat{\Lambda}_{NA}(t)$, we can also find another estimator of $S(t)$ (Fleming-Harrington):

$$\hat{S}_{FH}(t) = \exp(-\hat{\Lambda}_{NA}(t))$$

In general, this estimator of the survival function will be close to the Kaplan-Meier estimator, $\hat{S}_{KM}(t)$. We can also go the other way ... we can take the Kaplan-Meier estimate of $S(t)$, and use it to calculate an alternative estimate of the cumulative hazard function:

$$\hat{\Lambda}_{KM}(t) = -\log \hat{S}_{KM}(t)$$