# Chapter 4

# Comparison of Survival Curves

We spent the last class looking at some nonparametric approaches for estimating the survival function, $\hat{S}(t)$, over time for a single sample of individuals. Now we want to compare the survival estimates between two groups.

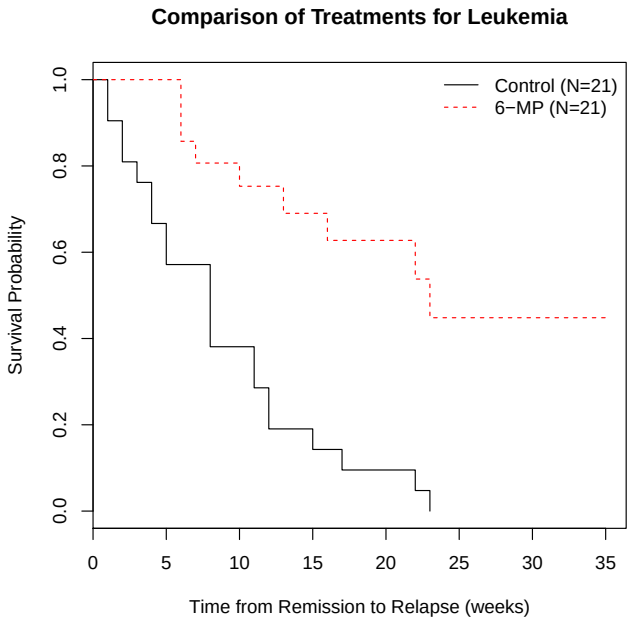


Figure 4.1: Time to remission of leukemia patients

**How can we form a basis for comparison?**

At a specific point in time, we could see whether the confidence intervals for the survival curves overlap.

However, the confidence intervals we have been calculating are **"pointwise"** $\Rightarrow$ they correspond to a confidence interval for $\hat{S}(t^*)$ at a single point in time, $t^*$. In other words, we can't say that the true survival function $S(t)$ is contained between the pointwise confidence intervals with 95% probability.

(**Aside:** if you're interested, the issue of confidence **bands** for the estimated survival function are discussed in Section 4.4 of Klein and Moeschberger)

Looking at whether the confidence intervals for $\hat{S}(t^*)$ overlap between the 6MP and placebo groups would only focus on comparing the two treatment groups at a single point in time, $t^*$.

**Should we base our overall comparison of $\hat{S}(t)$ on:**

- the furthest distance between the two curves?
- the median survival for each group?
- the average hazard? (for exponential distributions, this would be like comparing the mean event times)
- adding differences between the two survival estimates over time?

$$\sum_j \left[\hat{S}(t_{jA}) - \hat{S}(t_{jB})\right]$$

- a weighted sum of differences, where the weights reflect the number at risk at each time?
- a rank-based test? i.e., we could rank all of the event times, and then see whether the sum of ranks for one group was less than the other.

## 4.1 Nonparametric comparisons of groups

All of these are pretty reasonable options, and we'll see that there have been several proposals for how to compare the survival of two groups. For the moment, we are sticking to nonparametric comparisons.

**Why nonparametric?**

- **fairly robust**
- **efficient relative to parametric tests**
- **often simple and intuitive**

Before continuing the description of the two-sample comparison, I'm going to try to put this in a general framework to give a perspective of where we're heading in this class.

### 4.1.1 General Framework for Survival Analysis

We observe $(X_i, \delta_i, \mathbf{Z}_i)$ for individual $i$, where

- $X_i$ is a censored failure time random variable
- $\delta_i$ is the failure/censoring indicator
- $\mathbf{Z}_i$ represents a set of covariates

Note that $\mathbf{Z}_i$ might be a scalar (a single covariate, say treatment or gender) or may be a $(p \times 1)$ vector (representing several different covariates). These covariates might be:

- continuous
- discrete
- time-varying (more later)

If $\mathbf{Z}_i$ is a scalar and is binary, then we are comparing the survival of two groups, like in the leukemia example. More generally though, it is useful to build a model that characterizes the relationship between survival and all of the covariates of interest. We'll proceed as follows:

- Two group comparisons
- Multigroup and stratified comparisons - stratified logrank
- Failure time regression models
  - Cox proportional hazards model
  - Accelerated failure time model

## 4.2 Two sample tests

- Mantel-Haenszel logrank test
- Peto & Peto's version of the logrank test
- Gehan's Generalized Wilcoxon
- Peto & Peto's and Prentice's generalized Wilcoxon
- Tarone-Ware and Fleming-Harrington classes
- Cox's F-test (non-parametric version)

**References:**

| | |
|---|---|
| Collett | Section 2.5 |
| Klein & Moeschberger | Section 7.3 |
| Kleinbaum | Chapter 2 |
| Lee | Chapter 5 |

### 4.2.1 Mantel-Haenszel Logrank test

The logrank test is the most well known and widely used. It also has an intuitive appeal, building on standard methods for binary data. (Later we will see that it can also be obtained as the score test from a partial likelihood from the Cox Proportional Hazards model.)

First consider the following $(2 \times 2)$ table classifying those with and without the event of interest in a two group setting:

| | Event | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 0 | $d_0$ | $n_0 - d_0$ | $n_0$ |
| 1 | $d_1$ | $n_1 - d_1$ | $n_1$ |
| Total | $d$ | $n - d$ | $n$ |

If the margins of this table are considered fixed, then $d_0$ follows a *hypergeometric* distribution. Under the null hypothesis of no association between the event and group, it follows that

$$\begin{aligned} E(d_0) &= \frac{n_0 d}{n} \\ Var(d_0) &= \frac{n_0\, n_1\, d(n-d)}{n^2(n-1)} \end{aligned}$$

Therefore, under $H_0$:

$$\chi^2_{MH} = \frac{[d_0 - n_0\, d/n]^2}{\frac{n_0\, n_1\, d(n-d)}{n^2(n-1)}} \sim \chi^2_1$$

This is the Mantel-Haenszel statistic and is approximately equivalent to the Pearson $\chi^2$ test for equality of the two groups given by:

$$\chi^2_p = \sum \frac{(o-e)^2}{e}$$

Note: recall that the Pearson $\chi^2$ test was derived for the case where only the row margins were fixed, and thus the variance above was replaced by:

$$Var(d_0) = \frac{n_0\, n_1\, d(n-d)}{n^3}$$

### 4.2.2 Example:

Toxicity in a clinical trial with two treatments

| | Toxicity | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 0 | 8 | 42 | 50 |
| 1 | 2 | 48 | 50 |
| Total | 10 | 90 | 100 |

$$\chi^2_p = 4.00 \quad (p = 0.046)$$

$$\chi^2_{MH} = 3.96 \quad (p = 0.047)$$

Now suppose we have $K$ (2×2) tables, all independent, and we want to test for a common group effect. The Cochran-Mantel-Haenszel test for a common odds ratio not equal to 1 can be written as:

$$\chi^2_{CMH} = \frac{[\sum_{j=1}^{K}(d_{0j} - n_{0j} * d_j/n_j)]^2}{\sum_{j=1}^{K} n_{1j} n_{0j} d_j (n_j - d_j)/[n_j^2(n_j - 1)]}$$

where the subscript $j$ refers to the $j$-th table:

| | Event | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 0 | $d_{0j}$ | $n_{0j} - d_{0j}$ | $n_{0j}$ |
| 1 | $d_{1j}$ | $n_{1j} - d_{1j}$ | $n_{1j}$ |
| Total | $d_j$ | $n_j - d_j$ | $n_j$ |

This statistic is distributed approximately as $\chi^2_1$.

**How does this apply in survival analysis?**

Suppose we observe

**Group 1:** $(X_{11}, \delta_{11}) \ldots (X_{1n_1}, \delta_{1n_1})$

**Group 0:** $(X_{01}, \delta_{01}) \ldots (X_{0n_0}, \delta_{0n_0})$

We could just count the numbers of failures: eg., $d_1 = \sum_{j=1}^{K} \delta_{1j}$

### 4.2.3 Example: Leukemia data

Just counting up the number of remissions in each treatment group.

| | Fail | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 0 | 21 | 0 | 21 |
| 1 | 9 | 12 | 21 |
| Total | 30 | 12 | 42 |

$$\chi^2_p = 16.8 \quad (p = 0.001)$$
$$\chi^2_{MH} = 16.4 \quad (p = 0.001)$$

But, this doesn't account for the time at risk. Conceptually, we would like to compare the KM survival curves. Let's put the components side-by-side and compare.

### 4.2.4 Cox & Oakes Table 1.1 Leukemia example

| Ordered | Group 0 | | | Group 1 | | |
|---|---|---|---|---|---|---|
| Death Times | $d_j$ | $c_j$ | $r_j$ | $d_j$ | $c_j$ | $r_j$ |
| 1 | 2 | 0 | 21 | 0 | 0 | 21 |
| 2 | 2 | 0 | 19 | 0 | 0 | 21 |
| 3 | 1 | 0 | 17 | 0 | 0 | 21 |
| 4 | 2 | 0 | 16 | 0 | 0 | 21 |
| 5 | 2 | 0 | 14 | 0 | 0 | 21 |
| 6 | 0 | 0 | 12 | 3 | 1 | 21 |
| 7 | 0 | 0 | 12 | 1 | 0 | 17 |
| 8 | 4 | 0 | 12 | 0 | 0 | 16 |
| 9 | 0 | 0 | 8 | 0 | 1 | 16 |
| 10 | 0 | 0 | 8 | 1 | 1 | 15 |
| 11 | 2 | 0 | 8 | 0 | 1 | 13 |
| 12 | 2 | 0 | 6 | 0 | 0 | 12 |
| 13 | 0 | 0 | 4 | 1 | 0 | 12 |
| 15 | 1 | 0 | 4 | 0 | 0 | 11 |
| 16 | 0 | 0 | 3 | 1 | 0 | 11 |
| 17 | 1 | 0 | 3 | 0 | 1 | 10 |
| 19 | 0 | 0 | 2 | 0 | 1 | 9 |
| 20 | 0 | 0 | 2 | 0 | 1 | 8 |
| 22 | 1 | 0 | 2 | 1 | 0 | 7 |
| 23 | 1 | 0 | 1 | 1 | 0 | 6 |
| 25 | 0 | 0 | 0 | 0 | 1 | 5 |

We wrote down the number at risk for Group 1 for times 1-5 even though there were no events or censorings at those times.

## 4.3 Logrank Test: Formal Definition

The logrank test is obtained by constructing a $(2 \times 2)$ table at each distinct death time, and comparing the death rates between the two groups, conditional on the number at risk in the groups. The tables are then combined using the Cochran-Mantel-Haenszel test.

Note: The logrank is sometimes called the Cox-Mantel test. Let $t_1, ..., t_K$ represent the $K$ ordered, distinct death times. At the $j$-th death time, we have the following table:

| | Die/Fail | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 0 | $d_{0j}$ | $r_{0j} - d_{0j}$ | $r_{0j}$ |
| 1 | $d_{1j}$ | $r_{1j} - d_{1j}$ | $r_{1j}$ |
| Total | $d_j$ | $r_j - d_j$ | $r_j$ |

where $d_{0j}$ and $d_{1j}$ are the number of deaths in group 0 and 1, respectively at the $j$-th death time, and $r_{0j}$ and $r_{1j}$ are the number at risk at that time,

in groups 0 and 1. **The logrank test is:**

$$\chi^2_{logrank} = \frac{[\sum_{j=1}^{K}(d_{0j} - r_{0j} * d_j/r_j)]^2}{\sum_{j=1}^{K} \frac{r_{1j} r_{0j} d_j (r_j - d_j)}{[r_j^2 (r_j - 1)]}}$$

Assuming the tables are all independent, then this statistic will have an approximate $\chi^2$ distribution with 1 df.

**Based on the motivation for the logrank test, which of the survival-related quantities are we comparing at each time point?**

- $\sum_{j=1}^{K} w_j \left[\hat{S}_1(t_j) - \hat{S}_2(t_j)\right]$ **?**
- $\sum_{j=1}^{K} w_j \left[\hat{\lambda}_1(t_j) - \hat{\lambda}_2(t_j)\right]$ **?**
- $\sum_{j=1}^{K} w_j \left[\hat{\Lambda}_1(t_j) - \hat{\Lambda}_2(t_j)\right]$ **?**

**First several tables of leukemia data**

```
CMH analysis of leukemia data

TABLE 1 OF TRTMT BY REMISS
CONTROLLING FOR FAILTIME=1

TRTMT     REMISS

Frequency|
Expected |       0|       1|  Total
---------+--------+--------+
       0 |     19 |      2 |     21
         |     20 |      1 |
---------+--------+--------+
       1 |     21 |      0 |     21
         |     20 |      1 |
---------+--------+--------+
Total          40        2       42

TABLE 2 OF TRTMT BY REMISS
CONTROLLING FOR FAILTIME=2

TRTMT     REMISS

Frequency|
Expected |       0|       1|  Total
---------+--------+--------+
       0 |     17 |      2 |     19
         |  18.05 |   0.95 |
---------+--------+--------+
       1 |     21 |      0 |     21
         |  19.95 |   1.05 |
---------+--------+--------+
Total          38        2       40

TABLE 3 OF TRTMT BY REMISS
CONTROLLING FOR FAILTIME=3

TRTMT     REMISS

Frequency|
Expected |       0|       1|  Total
---------+--------+--------+
       0 |     16 |      1 |     17
         | 16.553 | 0.4474 |
---------+--------+--------+
       1 |     21 |      0 |     21
         | 20.447 | 0.5526 |
---------+--------+--------+
Total          37        1       38

TABLE 4 OF TRTMT BY REMISS
CONTROLLING FOR FAILTIME=4

TRTMT     REMISS

Frequency|
Expected |       0|       1|  Total
---------+--------+--------+
       0 |     14 |      2 |     16
         | 15.135 | 0.8649 |
---------+--------+--------+
       1 |     21 |      0 |     21
         | 19.865 | 1.1351 |
---------+--------+--------+
Total          35        2       37
```

**CMH statistic = logrank statistic**

```
               SUMMARY STATISTICS FOR TRTMT BY REMISS
                     CONTROLLING FOR FAILTIME


 Cochran-Mantel-Haenszel Statistics (Based on Table Scores)

Statistic   Alternative Hypothesis    DF       Value      Prob
------------------------------------------------------------------
    1        Nonzero Correlation        1      16.793     0.001
    2        Row Mean Scores Differ     1      16.793     0.001
    3        General Association        1      16.793     0.001 <===LOGRANK
                                                                 TEST
```

Note: Although CMH works to get the correct logrank test, it would require inputting the $d_j$ and $r_j$ at each time of death for each treatment group. There's an easier way to get the test statistic, which I'll show you shortly. **Calculating logrank statistic by hand: Leukemia Example:**

| Ordered | Group 0 | | Combined | | | | |
|---|---|---|---|---|---|---|---|
| Death Times | $d_{0j}$ | $r_{0j}$ | $d_j$ | $r_j$ | $e_j$ | $o_j - e_j$ | $v_j$ |
| 1 | 2 | 21 | 2 | 42 | 1.00 | 1.00 | 0.488 |
| 2 | 2 | 19 | 2 | 40 | 0.95 | 1.05 | |
| 3 | 1 | 17 | 1 | 38 | 0.45 | 0.55 | |
| 4 | 2 | 16 | 2 | 37 | 0.86 | 1.14 | |
| 5 | 2 | 14 | 2 | 35 | | | |
| 6 | 0 | 12 | 3 | 33 | | | |
| 7 | 0 | 12 | 1 | 29 | | | |
| 8 | 4 | 12 | 4 | 28 | | | |
| 10 | 0 | 8 | 1 | 23 | | | |
| 11 | 2 | 8 | 2 | 21 | | | |
| 12 | 2 | 6 | 2 | 18 | | | |
| 13 | 0 | 4 | 1 | 16 | | | |
| 15 | 1 | 4 | 1 | 15 | | | |
| 16 | 0 | 3 | 1 | 14 | | | |
| 17 | 1 | 3 | 1 | 13 | | | |
| 22 | 1 | 2 | 2 | 9 | | | |
| 23 | 1 | 1 | 2 | 7 | | | |
| Sum | | | | | | 10.251 | 6.257 |

In the previous table $o_j = d_{0j}$

$e_j = d_j r_{0j}/r_j$

$v_j = r_{1j} r_{0j} d_j (r_j - d_j)/[r_j^2 (r_j - 1)]$

$$\chi^2_{logrank} = \frac{(10.251)^2}{6.257} = 16.793$$

**Notes about logrank test:**

- The logrank statistic depends on ranks of event times only

- If there are no tied deaths, then the logrank has the form:

$$\frac{[\sum_{j=1}^{K}(d_{0j} - \frac{r_{0j}}{r_j})]^2}{\sum_{j=1}^{K} r_{1j}r_{0j}/r_j^2}$$

- Numerator can be interpreted as $\sum(o-e)$ where "o" is the observed number of deaths in group 0, and "e" is the expected number, given the risk set. The expected number equals #deaths $\times$ proportion in group 0 at risk.

- The $(o-e)$ terms in the numerator can be written as

$$\frac{r_{0j}r_{1j}}{r_j}(\hat{\lambda}_{1j} - \hat{\lambda}_{0j})$$

- It does not matter which group you choose to sum over. To see this, note that if we summed up (o-e) over the death times for the 6MP group we would get -10.251, and the sum of the variances is the same. So when we square the numerator, the test statistic is the same.

Analogous to the CMH test for a series of tables at different levels of a confounder, the logrank test is most powerful when "odds ratios" are constant over time intervals.

That is, it is most powerful for **proportional hazards**.

### 4.3.1 Checking the assumption of proportional hazards

- check to see if the estimated survival curves cross - if they do, then this is evidence that the hazards are <u>not</u> proportional

- more formal test: **any ideas?**

**What should be done if the hazards are <u>not proportional?</u>**

- If the difference between hazards has a consistent sign, the logrank test usually does well.

- Other tests are available that are more powerful against different alternatives.

### 4.3.2 Getting the logrank statistic using R:

```
               trt=Control
 time n.risk n.event survival std.err lower 95% CI upper 95% CI
    1     21       2   0.9048  0.0641      0.78754        1.000
    2     19       2   0.8095  0.0857      0.65785        0.996
    3     17       1   0.7619  0.0929      0.59988        0.968
    4     16       2   0.6667  0.1029      0.49268        0.902
    5     14       2   0.5714  0.1080      0.39455        0.828
    8     12       4   0.3810  0.1060      0.22085        0.657
   11      8       2   0.2857  0.0986      0.14529        0.562
   12      6       2   0.1905  0.0857      0.07887        0.460
   15      4       1   0.1429  0.0764      0.05011        0.407
   17      3       1   0.0952  0.0641      0.02549        0.356
   22      2       1   0.0476  0.0465      0.00703        0.322
   23      1       1   0.0000     NaN           NA           NA

               trt=6-MP
 time n.risk n.event survival std.err lower 95% CI upper 95% CI
    6     21       3    0.857  0.0764        0.720        1.000
    7     17       1    0.807  0.0869        0.653        0.996
   10     15       1    0.753  0.0963        0.586        0.968
   13     12       1    0.690  0.1068        0.510        0.935
   16     11       1    0.627  0.1141        0.439        0.896
   22      7       1    0.538  0.1282        0.337        0.858
   23      6       1    0.448  0.1346        0.249        0.807
```

The logrank test is given as follows:

```
             N Observed Expected (O-E)^2/E (O-E)^2/V
trt=Control 21       21     10.7      9.77      16.8
trt=6-MP    21        9     19.3      5.46      16.8

 Chisq= 16.8  on 1 degrees of freedom, p= 4.17e-05
```

## 4.4 Generalization of logrank test: Linear rank tests

The logrank and other tests can be derived by assigning scores to the ranks of the death times, and are members of a general class of **linear rank tests** (for more detail, see Lee, ch 5)

First, define

$$\hat{\Lambda}(t) = \sum_{j:t_j<t} \frac{d_j}{r_j}$$

where $d_j$ and $r_j$ are the number of deaths and the number at risk, respectively at the $j$-th ordered death time. Then assign these scores (suggested by Peto and Peto):

| Event | Score |
|---|---|
| Death at $t_j$ | $w_j = 1 - \hat{\Lambda}(t_j)$ |
| Censoring at $t_j$ | $w_j = -\hat{\Lambda}(t_j)$ |

To calculate the logrank test, simply sum up the scores for group 0.

### 4.4.1 Example

Group 0: 15, 18, 19, 19, 20

Group 1: 16+, 18+, 20+, 23, 24+

**Calculation of logrank as a linear rank statistic**

| Ordered Data | Group | $d_j$ | $r_j$ | $\hat{\Lambda}(t_j)$ | score $w_j$ |
|---|---|---|---|---|---|
| 15 | 0 | 1 | 10 | 0.100 | 0.900 |
| $16^+$ | 1 | 0 | 9 | 0.100 | -0.100 |
| 18 | 0 | 1 | 8 | 0.225 | 0.775 |
| $18^+$ | 1 | 0 | 7 | 0.225 | -0.225 |
| 19 | 0 | 2 | 6 | 0.558 | 0.442 |
| 20 | 0 | 1 | 4 | 0.808 | 0.192 |
| $20^+$ | 1 | 0 | 3 | 0.808 | -0.808 |
| 23 | 1 | 1 | 2 | 0.808 | 0.192 |
| $24^+$ | 1 | 0 | 1 | 0.808 | -0.808 |

The logrank statistic $S$ is sum of scores for group 0:

$$S = 0.900 + 0.775 + 0.442 + 0.442 + 0.192 = 2.75$$

The variance is:

$$Var(S) = \frac{n_0 n_1 \sum_{j=1}^{n} w_j^2}{n(n-1)}$$

In this case, $Var(S) = 1.210$, so

$$Z = \frac{2.75}{\sqrt{1.210}} = 2.50 \implies \chi^2_{logrank} = (2.50)^2 = 6.25$$

**Why is this form of the logrank equivalent?**

The logrank statistic S is equivalent to $\sum(o - e)$ over the distinct death times, where "$o$" is the observed number of deaths in group 0, and "$e$" is

the expected number, given the risk sets.

| | |
|---|---|
| **At deaths:** | weights are $1 - \hat{\Lambda}$ |
| **At censorings:** | weights are $-\hat{\Lambda}$ |

So we are summing up "1's" for deaths (to get $d_{0j}$), and subtracting $-\hat{\Lambda}$ at both deaths and censorings. This amounts to subtracting $d_j/r_j$ at each death or censoring time in group 0, at or after the $j$-th death. Since there are a total of $r_{0j}$ of these, we get $e = r_{0j} * d_j/r_j$.

### 4.4.2 Why is it called the *logrank* test?

Since $S(t) = \exp(-\Lambda(t))$, an alternative estimator of $S(t)$ is:

$$\hat{S}(t) = \exp(-\hat{\Lambda}(t)) = \exp(-\sum_{j:t_j<t} \frac{d_j}{r_j})$$

So, we can think of $\hat{\Lambda}(t) = -\log(\hat{S}(t))$ as yielding the "log-survival" scores used to calculate the statistic.

## 4.5 Comparing the CMH-type Logrank and "Linear Rank" logrank

### 4.5.1 CMH-type Logrank:

We motivated the logrank test through the CMH statistic for testing $H_o$ : $OR = 1$ over $K$ tables, where $K$ is the number of distinct death times. This turned out to be what we get when we use the logrank (default) option in Stata. (or the "STRATA" statement in SAS)

### 4.5.2 Linear Rank logrank

The linear rank version of the logrank test is based on adding up "scores" for one of the two treatment groups. The particular scores that gave us the same logrank statistic were based on the Nelson-Aalen estimator, i.e., $\hat{\Lambda} = \sum \hat{\lambda}(t_j)$. This is what you get when you use the "TEST" statement in SAS.

If there are no tied event times, then the two versions of the test will yield identical results. The more ties we have, the more it matters which version we use. The numerators of the two types of logrank tests will always be equivalent, but the denominators depend on the way ties are handled.

**CMH-type variance:**

$$\begin{aligned} var &= \sum \frac{r_{1j} r_{0j} d_j (r_j - d_j)}{r_j^2 (r_j - 1)} \\ &= \sum \frac{r_{1j} r_{0j}}{r_j (r_j - 1)} \frac{d_j (r_j - d_j)}{r_j} \end{aligned}$$

**Linear rank type variance:**

$$var = \frac{n_0 n_1 \sum_{j=1}^{n} w_j^2}{n(n-1)}$$

## 4.6 Gehan's Generalized Wilcoxon Test

First, let's review the Wilcoxon test for uncensored data:
Denote observations from two samples by:

$$(X_1, X_2, \ldots, X_n) \text{ and } (Y_1, Y_2, \ldots, Y_m)$$

Order the combined sample and define:

$$Z_{(1)} < Z_{(2)} < \cdots < Z_{(m+n)}$$

$$R_{i1} = \text{rank of } X_i$$

$$R_1 = \sum_{i=1}^{m+n} R_{i1}$$

Reject $H_0$ if $R_1$ is too big or too small, according to

$$\frac{R_1 - E(R_1)}{\sqrt{Var(R_1)}} \sim N(0,1)$$

where

$$\begin{aligned} E(R_1) &= \frac{m(m+n+1)}{2} \\ Var(R_1) &= \frac{mn(m+n+1)}{12} \end{aligned}$$

The **Mann-Whitney** form of the Wilcoxon is defined as:

$$U(X_i, Y_j) = U_{ij} = \begin{cases} +1 & \text{if} \quad X_i > Y_j \\ 0 & \text{if} \quad X_i = Y_j \\ -1 & \text{if} \quad X_i < Y_j \end{cases}$$

and

$$U = \sum_{i=1}^{n} \sum_{j=1}^{m} U_{ij}.$$

There is a simple correspondence between $U$ and $R_1$:

$$\begin{aligned} R_1 &= m(m+n+1)/2 + U/2 \\ \text{so} \qquad U &= 2R_1 - m(m+n+1) \end{aligned}$$

Therefore,

$$\begin{aligned} E(U) &= 0 \\ Var(U) &= mn(m+n+1)/3 \end{aligned}$$

### 4.6.1 Extending Wilcoxon to censored data

The Mann-Whitney form leads to a generalization for censored data. Define

$$U(X_i, Y_j) = U_{ij} = \begin{cases} +1 & \text{if} \quad x_i > y_j \ \text{ or } \ x_i^+ \geq y_j \\ 0 & \text{if} \quad x_i = y_i \ \text{ or lower value censored} \\ -1 & \text{if} \quad x_i < y_j \ \text{ or } \ x_i \leq y_j^+ \end{cases}$$

Then define

$$W = \sum_{i=1}^{n} \sum_{j=1}^{m} U_{ij}$$

Thus, there is a contribution to $W$ for every comparison where both observations are failures (except for ties), or where a censored observation is greater than or equal to a failure.

Looking at all possible pairs of individuals between the two treatment groups makes this a nightmare to compute by hand!

Gehan found an easier way to compute the above. First, pool the sample of $(n+m)$ observations into a single group, then compare each individual with the remaining $n+m-1$: For comparing the $i$-th individual with the $j$-th, define

$$U_{ij} = \begin{cases} +1 & \text{if} \quad t_i > t_j \ \text{ or } \ t_i^+ \geq t_j \\ -1 & \text{if} \quad t_i < t_j \ \text{ or } \ t_i \leq t_j^+ \\ 0 & \qquad\quad \textit{otherwise} \end{cases}$$

Then

$$U_i = \sum_{j=1}^{m+n} U_{ij}$$

Thus, for the $i$-th individual, $U_i$ is the number of observations which are definitely less than $t_i$ minus the number of observations that are definitely greater than $t_i$. We assume censorings occur after deaths, so that if $t_i = 18^+$ and $t_j = 18$, then we add 1 to $U_i$. The Gehan statistic is defined as

$$\begin{aligned} U &= \sum_{i=1}^{m+n} U_i \, \mathbf{1}_{\{i \text{ in group } 0\}} \\ &= W \end{aligned}$$

$U$ has mean 0 and variance

$$var(U) = \frac{mn}{(m+n)(m+n-1)} \sum_{i=1}^{m+n} U_i^2$$

### 4.6.2 Example from Lee

Group 0: 15, 18, 19, 19, 20

Group 1: 16+, 18+, 20+, 23, 24+

| Time | Group | $U_i$ | $U_i^2$ |
|---|---|---|---|
| 15 | 0 | -9 | 81 |
| $16^+$ | 1 | 1 | 1 |
| 18 | 0 | -6 | 36 |
| $18^+$ | 1 | 2 | 4 |
| 19 | 0 | -2 | 4 |
| 19 | 0 | -2 | 4 |
| 20 | 0 | 1 | 1 |
| $20^+$ | 1 | 5 | 25 |
| 23 | 1 | 4 | 16 |
| $24^+$ | 1 | 6 | 36 |
| SUM | | -18 | 208 |

$$U = -18$$

$$Var(U) = \frac{(5)(5)(208)}{(10)(9)} = 57.78$$

and

$$\begin{aligned} \chi^2 &= (-18)^2/57.78 \\ &= 5.61 \end{aligned}$$

**Example: (leukemia data)**

Use the sts test statement, with the appropriate option

```
             N Observed Expected (O-E)^2/E (O-E)^2/V
trt=Control 21    14.55     7.68      6.16      14.5
trt=6-MP    21     5.12    12.00      3.94      14.5

 Chisq= 14.5  on 1 degrees of freedom, p= 0.000143
```

## 4.7 Generalized Wilcoxon (Peto & Peto, Prentice)

Assign the following scores:

For a death at $t$: $\hat{S}(t+) + \hat{S}(t-) - 1$

For a censoring at $t$: $\hat{S}(t+) - 1$

The test statistic is $\sum(scores)$ for group 0.

| Time | Group | $d_j$ | $r_j$ | $\hat{S}(t+)$ | score $w_j$ |
|---|---|---|---|---|---|
| 15 | 0 | 1 | 10 | 0.900 | 0.900 |
| $16^+$ | 1 | 0 | 9 | 0.900 | -0.100 |
| 18 | 0 | 1 | 8 | 0.788 | 0.688 |
| $18^+$ | 1 | 0 | 7 | 0.788 | -0.212 |
| 19 | 0 | 2 | 6 | 0.525 | 0.313 |
| 20 | 0 | 1 | 4 | 0.394 | -0.081 |
| $20^+$ | 1 | 0 | 3 | 0.394 | -0.606 |
| 23 | 1 | 1 | 2 | 0.197 | -0.409 |
| $24^+$ | 1 | 0 | 1 | 0.197 | -0.803 |

$$\begin{aligned}\sum w_j \mathbf{1}_{\{j \text{ in group } 0\}} &= 0.900 + 0.688 + 2 * (0.313) + (-0.081) \\ &= 2.13\end{aligned}$$

$$\begin{aligned}Var(S) &= \frac{n_0 n_1 \sum_{j=1}^{n} w_j^2}{n(n-1)} = 0.765 \\ \text{so} \quad Z &= 2.13/0.765 = 2.433\end{aligned}$$

## 4.8 The Tarone-Ware class of tests:

This general class of tests is like the logrank test, but adds weights $w_j$. The logrank test, Wilcoxon test, and Peto-Prentice Wilcoxon are included as special cases.

$$\chi^2_{tw} = \frac{[\sum_{j=1}^{K} w_j(d_{1j} - r_{1j} * d_j/r_j)]^2}{\sum_{l=1}^{K} \frac{w_j^2 r_{1j} r_{0j} d_j (r_j - d_j)}{r_j^2 (r_j - 1)}}$$

| **Test** | **Weight** $w_j$ |
|---|---|
| Logrank | $w_j = 1$ |
| Gehan's Wilcoxon | $w_j = r_j$ |
| Peto/Prentice | $w_j = n\widehat{S}(t_j)$ |
| Fleming-Harrington | $w_j = [\hat{S}(t_j)]^\alpha$ |
| Tarone-Ware | $w_j = \sqrt{r_j}$ |

Note: these weights $w_j$ are not the same as the scores $w_j$ we've been talking about earlier, and they apply to the CMH-type form of the test statistic rather than $\sum(scores)$ over a single treatment group.

**Which test should we used?**

**CMH-type or Linear Rank?**
If there are not a high proportion of ties, then it doesn't really matter since:

- The two Wilcoxons are similar to each other
- The two logrank tests are similar to each other

Note: personally, I tend to use the CMH-type test, which you get with the STRATA statement in SAS and the TEST statement in STATA.

**Logrank or Wilcoxon?**

- Both tests have the right Type I power for testing the null hypothesis of equal survival, $H_o : S_1(t) = S_2(t)$
- The choice of which test may therefore depend on the alternative hypothesis, which will drive the power of the test.
- The Wilcoxon is sensitive to early differences between survival, while the logrank is sensitive to later ones. This can be seen by the relative weights they assign to the test statistic:

$$\begin{aligned} \text{LOGRANK} \quad \textit{numerator} &= \sum_j (o_j - e_j) \\ \text{WILCOXON} \quad \textit{numerator} &= \sum_j r_j (o_j - e_j) \end{aligned}$$

- The logrank is most powerful under the assumption of proportional hazards, which implies an alternative in terms of the survival functions of $H_a : S_1(t) = [S_2(t)]^\alpha$
- The Wilcoxon has high power when the failure times are lognormally distributed, with equal variance in both groups but a different mean. It will turn out that this is the assumption of an accelerated failure time model.
- Both tests will lack power if the survival curves (or hazards) "cross". However, that does not necessarily make them invalid!

## 4.9 $P$-sample and stratified logrank tests

We have been discussing two sample problems. In practice, more complex settings often arise:

- There are more than two treatments or groups, and the question of interest is whether the groups differ from each other.
- We are interested in a comparison between two groups, but we wish to adjust for another factor that may confound the analysis
- We want to adjust for lots of covariates.

We will first talk about comparing the survival distributions between more than 2 groups, and then about adjusting for other covariates.

### 4.9.1 $P$-sample logrank

Suppose we observe data from $P$ different groups, and the data from group $p$ ($p = 1, ..., P$) are:

$$(X_{p1}, \delta_{p1}) \ldots (X_{pn_p}, \delta_{pn_p})$$

We now contruct a $(P \times 2)$ table at each of the $K$ distinct death times, and compare the death rates between the $P$ groups, conditional on the number at risk.

Let $t_1, ....t_K$ represent the $K$ ordered, distinct death times.

At the $j$-th death time, we have the following table:

| | Die/Fail | | |
|---|---|---|---|
| Group | Yes | No | Total |
| 1 | $d_{1j}$ | $r_{1l} - d_{1j}$ | $r_{1j}$ |
| . | . | . | . |
| P | $d_{Pj}$ | $r_{Pj} - d_{Pj}$ | $r_{Pj}$ |
| Total | $d_j$ | $r_j - d_j$ | $r_j$ |

where $d_{pj}$ is the number of deaths in group $p$ at the $j$-th death time, and $r_{pj}$ is the number at risk at that time. The tables are then combined using the CMH approach.

If we were just focusing on this one table, then a $\chi^2_{(P-1)}$ test statistic could be constructed through a comparison of "o"s and "e"s, like before.

### 4.9.2 Example: Toxicity in a clinical trial with 3 treatments

```
TABLE OF GROUP BY TOXICITY

GROUP     TOXICITY

Frequency|
Row Pct  |no      |yes     |  Total
---------+--------+--------+
       1 |     42 |      8 |     50
         |  84.00 |  16.00 |
---------+--------+--------+
       2 |     48 |      2 |     50
         |  96.00 |   4.00 |
---------+--------+--------+
       3 |     38 |     12 |     50
         |  76.00 |  24.00 |
---------+--------+--------+
Total         128       22      150

STATISTICS FOR TABLE OF GROUP BY TOXICITY

Statistic                     DF     Value        Prob
------------------------------------------------------
Chi-Square                     2     8.097       0.017
Likelihood Ratio Chi-Square    2     9.196       0.010
Mantel-Haenszel Chi-Square     1     1.270       0.260

  Cochran-Mantel-Haenszel Statistics (Based on Table Scores)

Statistic   Alternative Hypothesis    DF    Value     Prob
-----------------------------------------------------------
   1        Nonzero Correlation        1    1.270    0.260
   2        Row Mean Scores Differ     2    8.043    0.018
   3        General Association        2    8.043    0.018
```

### 4.9.3 Formal Calculations:

Let $\mathbf{O}_j = (d_{1j}, ...d_{(P-1)j})^T$ be a vector of the observed number of failures in groups 1 to $(P-1)$, respectively, at the $j$-th death time. Given the risk sets $r_{1j}$, ... $r_{Pj}$, and the fact that there are $d_j$ deaths, then $\mathbf{O}_j$ has a distribution like a multivariate version of the hypergeometric. $\mathbf{O}_j$ has mean:

$$\mathbf{E}_j = (\frac{d_j\, r_{1j}}{r_j},\ \ldots\ , \frac{d_j\, r_{(P-1)j}}{r_j})^T$$

and variance covariance matrix:

$$\mathbf{V}_j = \begin{pmatrix} v_{11j} & v_{12j} & \ldots & v_{1(P-1)j} \\ & v_{22j} & \ldots & v_{2(P-1)j} \\ \ldots & & \ldots & \ldots \\ & & & v_{(P-1)(P-1)j} \end{pmatrix}$$

where the $\ell$-th diagonal element is:

$$v_{\ell\ell j} = r_{\ell j}(r_j - r_{\ell j})d_j(r_j - d_j)/[r_j^2(r_j - 1)]$$

and the $\ell m$-th off-diagonal element is:

$$v_{\ell m j} = r_{\ell j} r_{mj} d_j(r_j - d_j)/[r_j^2(r_j - 1)]$$

The resulting $\chi^2$ test for a single $(P \times 1)$ table would have (P-1) degrees and is constructed as follows:

$$(\mathbf{O}_j - \mathbf{E}_j)^T\, \mathbf{V}_j^{-1}\, (\mathbf{O}_j - \mathbf{E}_j)$$

### 4.9.4 Generalizing to K tables

Analogous to what we did for the two sample logrank, we replace the $\mathbf{O}_j$, $\mathbf{E}_j$ and $\mathbf{V}_j$ with the sums over the $K$ distinct death times. That is, let $\mathbf{O} = \sum_{j=1}^k \mathbf{O}_j$, $\mathbf{E} = \sum_{j=1}^k \mathbf{E}_j$, and $\mathbf{V} = \sum_{j=1}^k \mathbf{V}_j$. Then, the test statistic is:

$$(\mathbf{O} - \mathbf{E})^T\, \mathbf{V}^{-1}\, (\mathbf{O} - \mathbf{E})$$

**Example:**

Time taken to finish a test with 3 different noise distractions. All tests were stopped after 12 minutes.

| Noise Level | | |
|---|---|---|
| Group 1 | Group 2 | Group 3 |
| 9.0 | 10.0 | 12.0 |
| 9.5 | 12.0 | $12^+$ |
| 9.0 | $12^+$ | $12^+$ |
| 8.5 | 11.0 | $12^+$ |
| 10.0 | 12.0 | $12^+$ |
| 10.5 | 10.5 | $12^+$ |

**Observed data table**

| Ordered | Group 1 | | Group 2 | | Group 3 | | Combined | |
|---|---|---|---|---|---|---|---|---|
| Times | $d_{1j}$ | $r_{1j}$ | $d_{2j}$ | $r_{2j}$ | $d_{3j}$ | $r_{3j}$ | $d_j$ | $r_j$ |
| 8.5 | 1 | 6 | 0 | 6 | 0 | 6 | | |
| 9.0 | 2 | 5 | 0 | 6 | 0 | 6 | | |
| 9.5 | 1 | 3 | 0 | 6 | 0 | 6 | | |
| 10.0 | 1 | 2 | 1 | 6 | 0 | 6 | | |
| 10.5 | 1 | 1 | 1 | 5 | 0 | 6 | | |
| 11.0 | 0 | 0 | 1 | 4 | 0 | 6 | | |
| 12.0 | 0 | 0 | 2 | 3 | 1 | 6 | | |

**Expected table**

| Ordered | Group 1 | | Group 2 | | Group 3 | | Combined | |
|---|---|---|---|---|---|---|---|---|
| Times | $o_{1j}$ | $e_{1j}$ | $o_{2j}$ | $e_{2j}$ | $o_{3j}$ | $e_{3j}$ | $o_j$ | $e_j$ |
| 8.5 | | | | | | | | |
| 9.0 | | | | | | | | |
| 9.5 | | | | | | | | |
| 10.0 | | | | | | | | |
| 10.5 | | | | | | | | |
| 11.0 | | | | | | | | |
| 12.0 | | | | | | | | |

Doing the $P$-sample test by hand is cumbersome. Luckily, R and most other packages will do it for you! (or at least some version)

### $P$-sample logrank in R

The $P$-sample logrank test in R is as follows:

```
        N Observed Expected (O-E)^2/E (O-E)^2/V
group=1 6        6     1.57   12.4463   17.2379
group=2 6        5     4.53    0.0488    0.0876
group=3 6        1     5.90    4.0660    9.4495

 Chisq= 20.4  on 2 degrees of freedom, p= 3.75e-05
```

## 4.10 The Stratified Logrank

Sometimes, even though we are interested in comparing two groups (or maybe $P$) groups, we know there are other factors that also affect the outcome. It would be useful to adjust for these other factors in some way.

### 4.10.1 Example: The nursing home data

For the nursing home data, a logrank test comparing length of stay for those under and over 85 years of age suggests a significant difference (p=0.03).

However, we know that gender has a strong association with length of stay, and also age. Hence, it would be a good idea to STRATIFY the analysis by gender when trying to assess the age effect.

A **stratified logrank** allows one to compare groups, but allows the shapes of the hazards of the different groups to differ across strata. It makes the assumption that the group 1 vs group 2 hazard ratio is constant across strata. In other words: $\frac{\lambda_{1s}(t)}{\lambda_{2s}(t)} = \theta$ where $\theta$ is constant over the strata ($s = 1, ..., S$). This method of adjusting for other variables is not as flexible as that based on a modelling approach.

### 4.10.2 General setup for the stratified logrank

Suppose we want to assess the association between survival and a factor (call this $X$) that has two different levels. Suppose however, that we want to stratify by a second factor, that has $S$ different levels.

First, divide the data into $S$ separate groups. Within group $s$ ($s = 1, ..., S$), proceed as though you were constructing the logrank to assess the association between survival and the variable $X$. That is, let $t_{1s}, ..., t_{K_s s}$ represent the $K_s$ ordered, distinct death times in the $s$-th group.

At the $j$-th death time in group $s$, we have the following table:

| | Die/Fail | | |
|---|---|---|---|
| X | Yes | No | Total |
| 1 | $d_{s1j}$ | $r_{s1j} - d_{s1j}$ | $r_{s1j}$ |
| 2 | $d_{s2j}$ | $r_{s2j} - d_{s2j}$ | $r_{s2j}$ |
| Total | $d_{sj}$ | $r_{sj} - d_{sj}$ | $r_{sj}$ |

Let $O_s$ be the sum of the "o"s obtained by applying the logrank calculations in the usual way to the data from group $s$. Similarly, let $E_s$ be the sum of the "e"s, and $V_s$ be the sum of the "v"s.

The **stratified logrank** is

$$Z = \frac{\sum_{s=1}^{S}(O_s - E_s)}{\sqrt{\sum_{s=1}^{S}(V_s)}}$$

### 4.10.3 Stratified logrank using R

We run a stratified logrank test on the nursing home data. The output is as follows:

```
        <=85 >85
  Women  678 495
  Men    302 116

Asymptotic Two-Sample Logrank Test

data:  Surv(los, fail) by
 age85 (<=85, >85)
 stratified by gender
Z = -2.0556, p-value = 0.03982
alternative hypothesis: true theta is not equal to 1
```